

Poisson Regression

STA 310: Homework 4 - Amaris

Instructions:

- Write all narrative using full sentences. Write all interpretations and conclusions in the context of the data.
- Be sure all analysis code is displayed in the rendered pdf.
- If you are fitting a model, display the model output in a neatly formatted table. (The tidy and kable functions can help!)
- If you are creating a plot, use clear and informative labels and titles.
- Make sure to upload to both Gradescope and Canvas in a reproducible format per the instructions of prior homework assignments.

These exercises are derived from BMLR, Chapter 4.

1.

Answer parts a - d in the context of the following study:

A state wildlife biologist collected data from 250 park visitors as they left at the end of their stay. Each was asked to report the number of fish they caught during their one-week stay. On average, visitors caught 21.5 fish per week.

a. **Define the response.**

The response variable is the number of fish caught by a park visitor during their one-week stay.

b. **What are the possible values for the response?**

The response takes non-negative integer values (0,1,2,3,...) since fish counts cannot be negative or fractional.

c. **What does λ represent?**

In a Poisson model, λ represents the expected (mean) number of fish caught per visitor per week. Here, based on the sample, $\lambda = 21.5$.

d. **Would a zero-inflated model be considered here? If so, what would be a “true zero”?**

Yes, a zero-inflated model could be considered if there are more zeros than expected under a standard Poisson model. A true zero would correspond to visitors who did not fish at all during their stay (structural zeros), rather than visitors who fished but happened to catch zero fish by chance.

2.

@brockmann1996 carried out a study of nesting female horseshoe crabs. Female horseshoe crabs often have male crabs attached to a female's nest known as satellites. One objective of the study was to determine which characteristics of the female were associated with the number of satellites. Of particular interest is the relationship between the width of the female carapace and satellites.

The data can be found in `crab.csv` in the data folder. It includes the following variables:

- Satellite = number of satellites
- Width = carapace width (cm)
- Weight = weight (kg)
- Spine = spine condition (1 = both good, 2 = one worn or broken, 3 = both worn or broken)
- Color = color (1 = light medium, 2 = medium, 3 = dark medium, 4 = dark)

Make sure to convert Spine and Color to the appropriate data types in R before doing the analysis.

```
crab <- read.csv("crab.csv")
library(ggplot2)
library(broom)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'


## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(knitr)
```

- a. Create a histogram of Satellite. Is there preliminary evidence the number of satellites could be modeled as a Poisson response? Briefly explain.

```
ggplot(crab, aes(x = Satellite)) +
  geom_histogram(binwidth = 1, color = "white") +
  labs(
    title = "Histogram of Number of Satellites",
    x = "Number of Satellites",
    y = "Count"
  ) +
  theme_minimal()
```



Yes — there is preliminary evidence that the number of satellites could be modeled as a Poisson response.

The histogram shows that Satellite is a count variable taking non-negative integer values, with a right-skewed distribution and many small counts, which is consistent with a Poisson-type process. However, there is also a large spike at zero and a long right tail, suggesting possible overdispersion or excess zeros, so while a Poisson model is a reasonable starting point, it may not fully capture the variability in the data.

- b. Fit a Poisson regression model including Width, Weight, and Spine as predictors. Display the model with the 95% confidence interval for each coefficient.

```
# Fit Poisson regression
crab_pois <- glm(
  Satellite ~ Width + Weight + Spine,
  family = poisson(link = "log"),
  data = crab
)

# Display model summary
crab_pois %>%
  tidy(conf.int = TRUE, conf.level = 0.95) %>%
  mutate(across(where(is.numeric), round, 3)) %>%
  kable( caption = "Poisson Regression of Number of Satellites on Female Characteristics",
         col.names = c("Term",
                       "Estimate",
                       "Std. Error",
                       "z value",
                       "p-value",
                       "Lower 95% CI", "Upper 95% CI"
```

```
)
)
```

```
## Warning: There was 1 warning in 'mutate()'.
## i In argument: 'across(where(is.numeric), round, 3)'.
## Caused by warning:
## ! The '...' argument of 'across()' is deprecated as of dplyr 1.1.0.
## Supply arguments directly to '.fns' through an anonymous function instead.
##
## # Previously
## across(a:b, mean, na.rm = TRUE)
##
## # Now
## across(a:b, \(x) mean(x, na.rm = TRUE))
```

Table 1: Poisson Regression of Number of Satellites on Female Characteristics

Term	Estimate	Std. Error	z value	p-value	Lower 95% CI	Upper 95% CI
(Intercept)	-1.263	0.905	-1.396	0.163	-3.033	0.512
Width	0.047	0.047	1.008	0.313	-0.045	0.138
Weight	0.000	0.000	2.705	0.007	0.000	0.001
Spine	-0.016	0.054	-0.295	0.768	-0.119	0.091

c. **Describe the effect of Spine in terms of the mean number of satellites.**

Because this is a Poisson regression with a log link, the coefficient for Spine represents a multiplicative effect on the mean number of satellites. Holding width and weight constant, a one-unit increase in spine condition (i.e., worse spine condition) is associated with a factor of $\exp(-0.016) = 0.984$ change in the expected number of satellites, corresponding to about a 1.6% decrease in the mean number of satellites. However, this effect is not statistically significant (95% CI includes 0; $p=0.768$), indicating little evidence that spine condition is meaningfully associated with satellite count.

3. Use the scenario from the previous exercise to answer questions (a) - (d).
 - a. We would like to fit a quasi-Poisson regression model for this data. Briefly explain why we may want to consider fitting a quasi-Poisson regression model for this data.
 - b. Fit a quasi-Poisson regression model that corresponds with the model chosen in the previous exercise. Display the model.
 - c. What is the estimated dispersion parameter? Show how this value is calculated.
 - d. How do the estimated coefficients change compared to the model chosen in the previous exercise? How do the standard errors change?
4. The goal of this exercise is to use simulation to understand the equivalency between a gamma-Poisson mixture and a negative binomial distribution.

Remember to set a seed so your simulations are reproducible!

- a. Use the R function `rpois()` to generate 10,000 x_i from a regular Poisson distribution, $X \sim \text{Poisson}(\lambda = 1.5)$. Plot a histogram of this distribution and note its mean and variance. Next, let $Y \sim \text{Gamma}(r = 3, \lambda = 2)$ and use `rgamma()` to generate 10,000 random y_i from this distribution.

Now, consider 10,000 different Poisson distributions where $\lambda_i = y_i$. Randomly generate one z_i from each Poisson distribution. Plot a histogram of these z_i and compare it to your original histogram of X (where $X \sim \text{Poisson}(1.5)$). How do the means and variances compare?

- b. A negative binomial distribution can actually be expressed as a gamma-Poisson mixture. In Part a, you looked at a gamma-Poisson mixture $Z \sim \text{Poisson}(\lambda)$ where $\lambda \sim \text{Gamma}(r = 3, \lambda' = 2)$.

Find the parameters of a negative binomial distribution $X \sim \text{NegBinom}(r, p)$ such that X is equivalent to Z . As a hint, the means of both distributions must be the same, so $\frac{r(1-p)}{p} = \frac{3}{2}$.

Show through histograms and summary statistics that your negative binomial distribution is equivalent to the gamma-Poisson mixture. You can use `rnbinom()` in R.

- c. Make an argument that if you want a $\text{NegBinom}(r, p)$ random variable, you can instead sample from a Poisson distribution, where the λ values are themselves sampled from a gamma distribution with parameters r and $\lambda' = \frac{p}{1-p}$. You may show equivalency via the simulations or mathematically, however, make sure your arguments are precise and clear.

5. In a 2018 study, Chapp et al. (2018) scraped every issue statement from webpages of candidates for the U.S. House of Representatives, counting the number of issues candidates commented on and scoring the level of ambiguity of each statement. We will focus on the issue counts, and determining which attributes (of both the district as a whole and the candidates themselves) are associated with candidate silence (commenting on 0 issues) and a willingness to comment on a greater number of issues. The data set is in `ambiguity.csv` in the data folder. This analysis will focus on the following variables:

- `name` : candidate name
- `distID` : unique identification number for Congressional district
- `ideology` : candidate left-right orientation
- `democrat` : 1 if Democrat, 0 if Republican
- `totalIssuePages` : number of issues candidates commented on (response)

See @roback2021beyond for the full list of variables.

We will use a **hurdle model** to analyze the data. A hurdle model is similar to a zero-inflated Poisson model, but instead of assuming that “zeros” are comprised of two distinct groups—those who would always be 0 and those who happen to be 0 on this occasion—the hurdle model assumes that “zeros” are a single entity. Therefore, in a hurdle model, cases are classified as either “zeros” or “non-zeros”, where “non-zeros” *hurdle* the 0 threshold—they must always have counts of 1 or above.

We will use the `pscl` package and the `hurdle` function in it to analyze a hurdle model. Note that coefficients in the “zero hurdle model” section of the output relate predictors to the log-odds of being a *non-zero* (i.e., having at least one issue statement), which is opposite of the ZIP model.

- a. Visualize the distribution of the response variable `totalIssuePages`. Why might we consider using a hurdle model compared to a Poisson model? Why is a zero-inflated Poisson model not appropriate in this scenario?

- b. Create a plot of the empirical log odds of having at least one issue statement by ideology. You may want to group ideology values first. What can you conclude from this plot?
- c. Create a hurdle model with ideology and democrat as predictors in both parts. Display the model. Interpret ideology in both parts of the model.
- d. Repeat (d), but include an interaction in both parts. Interpret the interaction in the zero hurdle part of the model.

Grading

Total	39
Ex 1	4
Ex 2	6
Ex 3	8
Ex 4	8
Ex 5	10
Workflow & formatting	3

The “Workflow & formatting” grade is to based on the organization of the assignment write up along with the reproducible workflow. This includes having an organized write up with neat and readable headers, code, and narrative, including properly rendered mathematical notation. It also includes having a reproducible .Rmd document that can be rendered to reproduce the submitted PDF.

Extra resources and hints for Exercise 5 (Hurdle Problem)

1. There is an article on hurdle models that has helped some students that can be found here:

<https://jsdajournal.springeropen.com/articles/10.1186/s40488-021-00121-4>

2. There are two parts to the hurdle model, the count part and the binary part of the model and how to run it can be found in the R documentation.
3. You are not required to account for overdispersion (or check for it), however, if you do, this is great and just please make sure to think through how to do this properly as it involves integrating multiple parts of the Poisson lectures.
4. How do we handle the data in 5b? Do we bin it?

Yes, you should bin it or group it using the function `cut_interval`.

For example, something like this might help:

`ideologybin = cut_interval(ideology, n=5)`, where the number of bins of 5 was chosen empirically by playing around with the data.

References